

Week 7

???

7.1 Zorn's Lemma and Intro to Modules Over PIDs

2/13:

- Picking up from last time with Zorn's lemma.
- **Partially ordered set:** A set together with a binary relation indicating that, for certain pairs of elements in the set, one of the elements precedes the other in the ordering. *Also known as poset. Denoted by P .*
 - The domain of the **partial order** may be a proper subset of $P \times P$.
- **Partial order:** The binary relation on a poset.
- **Maximal** ($f \in P$): An element $f \in P$ such that for all $q \in P$, the statement $q > f$ is false.
- Example.
 - Let X be a set with $|X| \geq 2^{[1]}$.
 - Define a poset $P = \{A \subseteq X\}$ with corresponding partial order defined by taking subsets. In particular, if $A \subset B$, write $A \leq B$.
 - For any $x \in X$, $X - \{x\}$ is a maximal element of P .
- **Chain:** A subset of a poset P such that if c_1, c_2 are in said subset, then implies $c_1 \leq c_2$ or $c_2 \leq c_1$. *Denoted by C .*
 - In other words, a chain is a subset of a poset that is a **totally ordered set**.
- **Totally ordered set:** A set together with a binary relation indicating that, for any pair of elements in the set, one of the elements precedes the other in the ordering.
- Observation: If F is a subset of a nonempty finite chain C , then there exists $c \in F$ such that $c \geq q$ for all $q \in F$.
- **Upper bound** (of C): An element $p \in P$ such that $p \geq c$ for all $c \in C$.
- **Zorn's lemma:** Let P be a poset that satisfies
 - (i) $P \neq \emptyset$;
 - (ii) Every chain $C \subset P$ has an upper bound.

Then P has a maximal element.

¹Nori denotes cardinality by $\#X$.

- We will not prove Zorn's lemma. It rarely if ever gets proven in an undergraduate course, maybe in a logic course.
 - And by “prove” we mean “deduce Zorn's lemma from the Axiom of Choice.”
- We now investigate a situation in which Zorn's lemma gets applied.
- Let M be a finitely generated A -module.
 - Let $v_1, \dots, v_r \in M$ be elements such that $M = Av_1 + \dots + Av_r$.
 - Before we prove the proposition that requires Zorn's lemma, we will need one more definition: that of a **maximal submodule**.
- **Maximal submodule** (of M): A submodule of M that is a maximal element of the poset

$$P = \{N \subsetneq M : N \text{ is an } A\text{-submodule}\}$$

- Proposition: Every nonzero finitely generated A -module M has a maximal submodule.

Proof. To prove that M has a maximal submodule, it will suffice show that there exists a maximal element of the poset

$$P = \{N \subsetneq M : N \text{ is an } A\text{-submodule}\}$$

To do this, Zorn's lemma tells us that it will suffice to confirm that $P \neq \emptyset$ and that every chain $C \subset P$ has an upper bound. Let's begin.

We first confirm that $P \neq \emptyset$. By hypothesis, M is nonzero. Thus, the zero A -submodule is a proper subset of M , so $0 \in P$ and hence P is nonempty.

We now confirm that every chain $C \subset P$ has an upper bound. Let $C \subset P$ be an arbitrary chain. Define

$$\mathcal{N}_C = \bigcup \{N : N \in C\}$$

We will first verify that $\mathcal{N}_C \in P$, and then we will show that \mathcal{N}_C is an upper bound of C . Let's begin. To verify that $\mathcal{N}_C \in P$, it will suffice to demonstrate that \mathcal{N}_C is an A -submodule of M and that $\mathcal{N}_C \subsetneq M$.

To demonstrate that \mathcal{N}_C is an A -submodule, Proposition 10.1 tells us that it will suffice to show that $\mathcal{N}_C \neq \emptyset$ and $n_1 + an_2 \in \mathcal{N}_C$ for all $a \in A$ and $n_1, n_2 \in \mathcal{N}_C$. Since P is nonempty, \mathcal{N}_C is nonempty by definition, as desired. Additionally, let $n_1, n_2 \in \mathcal{N}_C$ be arbitrary. It follows by the definition of \mathcal{N}_C that there exist $N_1, N_2 \in C$ such that $n_i \in N_i$ ($i = 1, 2$). WLOG, assume $N_1 \subset N_2$. Then $n_1, n_2 \in N_2$. It follows since N_2 is an A -submodule that $n_1 + an_2 \in N_2 \subset \mathcal{N}_C$ for all $a \in A$, as desired.

We know that $\mathcal{N}_C \subset M$. Thus, if $\mathcal{N}_C \subsetneq M$, then we must have $\mathcal{N}_C = M$. Suppose for the sake of contradiction that $\mathcal{N}_C = M$. Recall that $M = Av_1 + \dots + Av_r$. Since the v_i are elements of M and $\mathcal{N}_C = M$, it follows that $v_i \in \mathcal{N}_C$ ($i = 1, \dots, r$). Thus, as before, there must exist $N_1, \dots, N_r \in C$, not necessarily distinct, such that $v_i \in N_i$ ($i = 1, \dots, r$). It follows by the observation from earlier that there is an $i \in [r]$ such that for all $j \in [r]$, $N_j \subset N_i$. Consequently, $v_j \in N_j \subset N_i$ ($j = 1, \dots, r$). But N_i is an A -submodule, so $M = Av_1 + \dots + Av_r \subset N_i \subset M$. But this means that $N_i = M$, contradicting the assumption that $N_i \subsetneq P$ (since $N_i \in P$). Therefore, $\mathcal{N}_C \subsetneq M$, as desired.

It follows that $\mathcal{N}_C \in P$, as desired. Lastly, we have by its definition that $N \subset \mathcal{N}_C$ for all $N \in C$, meaning that \mathcal{N}_C is an upper bound of C by definition. Therefore, by Zorn's lemma, P has a maximal element, and hence M has a maximal submodule, as desired. \square

- Corollary: Every nonzero commutative ring R has a maximal ideal.

Proof. Consider R as an R -module. Then $R = (1)$ is finitely generated. This combined with the fact that it is nonzero by hypothesis allows us to invoke the above proposition, learning that R has a maximal submodule N . But by the observation from Lecture 6.1, N is a left ideal, which is equivalent to a two-sided ideal in a commutative ring. Maximality transfers over as well (as we can confirm), proving that N is the desired maximal ideal of R . \square

- Remark: Suppose that J is a two-sided ideal of A . Let M be an A -module such that for all $a \in J$ and $m \in M$, we have $am = 0$. Then M may be regarded as an (A/J) -module in a natural manner.
 - In particular, we may take $\rho : A \rightarrow \text{End}(M, +)$ to be a ring homomorphism.
 - We can factor $\rho = \bar{\rho} \circ \pi$, where $\pi : A \rightarrow A/J$ and $\bar{\rho} : A/J \rightarrow \text{End}(M, +)$. It follows that $\bar{\rho}$ is a ring homomorphism. Therefore, M is an A/J -module.
 - This remark will be used!
 - Review annihilators from Section 10.1!
- Remark: Given a left ideal $I \subset A$ and an A -module M , we get a whole lot of modules because each element of M generates one. In particular, we note that $Im \subset Am \subset M$, where both Im, Am are submodules for all $m \in M$.

- **Product** (of modules): The A -submodule of M defined as follows. Denoted by IM . Given by

$$IM = \sum_{m \in M} Im$$

- It follows that M/IM is an A -module, but also one with a special property: $a(M/IM) = 0$ for all $a \in I$.
 - If A is commutative, then M/IM is an A/I -module.
- Proposition: Let R be a nonzero commutative ring. If $R^m \cong R^n$ as R -modules, then $m = n$.

Proof. Let $I \subset R$ be a maximal ideal. (We know that one exists by the above corollary.) If $f : R^m \rightarrow R^n$ is an isomorphism of R -modules, then f restricts to $I(R^m) \rightarrow I(R^n)$. This gives rise to the isomorphism $\bar{f} : R^m/I(R^m) \rightarrow R^n/I(R^n)$ of R -modules, in fact of R/I modules. It follows that R/I is a field, so $m = n$. \square

- Classifying modules up to isomorphism under commutative rings.
 - This is a hard problem, and there are still many open problems in this field today.
 - We will not go into this, though.
- We now move on to modules over PIDs.
 - Nori will go *much* slower than the book.
 - Do you have any recommended resources??
 - Do we need to read and understand Chapters 10-11 to start on Chapter 12??
- Objective: Let R be a PID. Classify all finitely generated R -modules up to isomorphism.
 - Our first result in this field was that submodules of R^n are equal to R^m for $m \leq n$.
 - Where this is applicable: \mathbb{Z} and $F[X]$.
 - Go back and check out \mathbb{Z} -modules and $F[X]$ -modules in Section 10.1!
- **Torsion module:** An R -module M such that for all $m \in M$, there exists $0 \neq a \in R$ such that $am = 0$.
- **Torsion-free module:** An R -module M such that for all nonzero $m \in M$ and for all nonzero $a \in R$, we have $am \neq 0$.
- Theorem: If M is a finitely generated torsion-free R -module, then $M \cong R^n$ for some n .
 - With a little work, we could prove this. But Nori will postpone it.

- **p -primary** (module): An R -module M such that for all $m \in M$, there exists $k \geq 0$ for which $p^k m = 0$, where p is prime in R .
- We want to classify these up to isomorphism.
 - Nori can state these today, but will not have time to prove it until another day.
 - Something that gets annihilated by p is a $\mathbb{Z}/(p)$ -module. The moment you go from $k = 1$ to $k = 2$, things get interesting.
- Examples: $R/(p^{n_1}) \oplus \cdots \oplus R/(p^{n_k})$, where $n_1 \geq \cdots \geq n_k \geq 1$.
 - Note that $k = 0$ is allowed.
- Uniqueness will take some time, but existence can be given as an exercise now.
- M/pM is an $R/(p)$ -vector space. pM/p^2M is an $R/(p)$ -vector space as well. So is $p^k M/p^{k+1}M$.
 - Use d_0, d_1, \dots, d_k to denote the dimensions of the vector spaces.
 - d_0, \dots, d_k is a decreasing sequence of nonnegative integers.

7.2 Office Hours (Nori)

- Homework questions.
 - See pictures + unnumbered lemma.
 - Example of the kernel being bigger than (f) .
 - A ring homomorphism $\mathbb{Z}[X] \rightarrow \mathbb{R}$ must be evaluation by the universal property of polynomial rings.
 - Factoring enables a constraint on a .
- Lecture 6.1: Proposition proof?
- Lecture 6.1: $(2) \subsetneq \mathbb{Z}$ example?
- Lecture 6.1: The end of the theorem proof.
- Lecture 6.2: Does the first theorem you proved not appear in the book until Chapter 12?
- Lecture 6.2: What is A in the proof?
- Resources for the proofs in Week 6?
- Lecture 7.1: Quotient stuff.
- Recommended resources for modules over PIDs? Chapter 12?
 - We should be able to read chapter 12, since chapter 11 is just vector spaces.
 - Nori's doing Chapter 12 in the classical manner (pre-1970). Dummit and Foote (2004) just does it in the first few pages as the **elementary divisor theorem**.
- HW6: So you want us to solve 1, 10, 13 for our own edification, but we don't need to write up a solution? Will we ever be responsible for the content therein?
 - We'll need to understand them to move forward.
 - Q6.4-Q6.5 are particularly important (good for number theory).

7.3 Office Hours (Ray)

- Universal properties save you from having to do pages upon pages of ring homomorphism checks (think Q3.10).
- Algebra: Chapter 0 by Paolo Aluffi for learning quotienting by polynomials.
 - Universal properties show up on page 30.
 - Read stuff before as needed.
 - Has a chapter called universal properties of polynomial rings. Universal properties of quotients, too.
- Direct sums and direct products.
 - Let M, N be R -modules. Then $M \times N$ is an R -module defined by the Cartesian product of the sets and with **diagonal** module action $r(m, n) = (rm, rn)$ (diagonal meaning we just act on two elements).
 - $M \oplus N = M \times N$.
 - For infinite sets, we get a difference. Indeed, $\prod_{i=1}^{\infty} M_i \neq \bigoplus_{i=1}^{\infty} M_i$.

7.4 Classifying Modules Over PIDs

- 2/15:
- We pick up from yesterday, classifying finitely generated R -modules M up to isomorphism when R is a PID.
 - In particular, we begin with a further investigation of the properties of torsion modules.
 - **Lift** (of $x \in M/M'$): The choice of an element $y \in M$ such that $\pi(y) = x$.
 - Lemma:
 - (i) $\text{Tor}(M)$ is an R -submodule of M .

Proof. To prove that $\text{Tor}(M)$ is an R -submodule of M , Proposition 10.1 tells us that it will suffice to show that $\text{Tor}(M) \neq \emptyset$ and that $x + ry \in \text{Tor}(M)$ for all $r \in R$, $x, y \in \text{Tor}(M)$. Consider $0 \in M$. By definition, $r \cdot 0 = 0$. Thus, $0 \in \text{Tor}(M)$ as desired. Additionally, let $r \in R$ and $x, y \in \text{Tor}(M)$ be arbitrary. Since $x, y \in \text{Tor}(M)$, there exist nonzero $a, b \in R$ such that $ax = 0$ and $by = 0$. Because R is an integral domain (as a PID), a, b nonzero implies that $ab \neq 0$. Thus, since

$$ab(x + ry) = abx + abry = b(ax) + ar(by) = b(0) + ar(0) = 0$$

we have that $x + ry \in \text{Tor}(M)$, as desired. \square

- (ii) The quotient module $M/\text{Tor}(M)$ is torsion-free.

Proof. To prove that $M/\text{Tor}(M)$ is torsion-free, it will suffice to show that every torsion element of $M/\text{Tor}(M)$ is 0. Let's begin. Let $v \in M/\text{Tor}(M)$ be an arbitrary torsion element. Then there exists $a \in R$ nonzero such that $av = 0$. Now lift $v \in M/\text{Tor}(M)$ to $w \in M$. The constraint $av = 0 = 0 + \text{Tor}(M)$ from the quotient module implies that $0 = a\pi(w) = \pi(aw)$, hence $aw \in \text{Tor}(M)$. Thus, there exists $b \in R$ nonzero such that $b(aw) = 0$. It follows that $(ba)w = 0$, where $ba \neq 0$ since $a, b \neq 0$ by the fact that R is an integral domain. Thus, $w \in \text{Tor}(M)$, and hence $v = \pi(w) = 0$, as desired. \square

- We now give some claims that will be useful later today, but whose proofs we will delay until next lecture.
- The first one pertains to the properties of finitely generated torsion-free modules over an integral domain.

- Lemma: Let R be an integral domain, and let M be a finitely generated R -module. Then there exists a submodule $M' \subset M$ such that...
 - (i) $M' \cong R^h$ for some $h \geq 0$;
 - (ii) There exists a nonzero $a \in R$ such that $aM \subset M'$ (equivalently, $a(M/M') = 0$).
- The next two pertain to the properties of finitely generated modules over a PID.
- Corollary: Every finitely generated torsion-free module M over a PID R is isomorphic to R^h for some $h \in \mathbb{Z}_{\geq 0}$.
- Theorem: Let M be a finitely generated R -module, where R is a PID. Then...
 - (i) $\text{Tor}(M) \oplus R^h \cong M$ for some $h \geq 0$;
 - (ii) $\text{Tor}(M)$ is finitely generated.
- **Rank** (of a module): The number h pertaining to an R -module M , where $M/\text{Tor}(M) \cong R^h$. Denoted by **rank**(M).
 - It follows by the proposition from last lecture (Lecture 7.1) that rank is well-defined.
- Corollary: Finitely generated R -modules M_1 and M_2 are isomorphic to each other iff
 - (i) M_1 and M_2 have the same rank;
 - (ii) $\text{Tor}(M_1)$ is isomorphic to $\text{Tor}(M_2)$.

Proof. Suppose first that $\phi : M_1 \rightarrow M_2$ is an isomorphism. Then naturally they will have the same ranks and torsion submodules.

On the other hand, if $\text{rank}(M_1) = \text{rank}(M_2)$, then $M_1/\text{Tor}(M_1) \cong M_2/\text{Tor}(M_2)$. This combined with the hypothesis that $\text{Tor}(M_1) \cong \text{Tor}(M_2)$ implies that

$$\begin{aligned} \text{Tor}(M_1) \oplus M_1/\text{Tor}(M_1) &\cong \text{Tor}(M_2) \oplus M_2/\text{Tor}(M_2) \\ M_1 &\cong M_2 \end{aligned}$$

where the second line follows from the preceding theorem. □

- The classification of finitely generated R -modules (R a PID) is completed by the following results.
- **p -primary component** (of a module): The submodule of a module M consisting of those $m \in M$ such that $p^k m = 0$ for some $k \in \mathbb{Z}_{\geq 0}$. Denoted by $M_{(p)}$.
 - Showing that $M_{(p)}$ is a submodule of M can be accomplished with the submodule criterion (Proposition 10.1), just like in the first lemma proven today.
- Notation and observations.
 1. Let M_1, \dots, M_k be submodules of M . Then $T : \prod_{i=1}^k M_i \rightarrow M$ defined by

$$T(m_1, \dots, m_k) = m_1 + \dots + m_k$$
 is not injective in general.
 - For example, if $k = 2$, then $\ker(T) \cong M_1 \cap M_2$ in general.
 - Thus, some care is required in our selection of submodules if we want $\ker(T) = 0$.
 2. Obtaining a natural R -module homomorphism $T : \oplus_{i \in I} M_i \rightarrow M$ defined as above.
 - We have that $\oplus_{i \in I} M_i \subset \prod_{i \in I} M_i$ in general. Here's why:
 - Given a finite subset $F \subset I$, we may regard $\prod_{i \in F} M_i$ as a submodule of $\prod_{i \in I} M_i$ by taking the entries in the i^{th} place to be zero for all $i \notin F$.

- The direct sum is simply the union of the submodules $\prod_{i \in F} M_i$ taken over all finite $F \subset I$.
- We define T on the overall direct sum one submodule $\prod_{i \in F} M_i$ at a time.
- Proposition: The natural R -module homomorphism $T : \bigoplus_{(p)} M_{(p)} \rightarrow \text{Tor}(M)$ is an isomorphism, where the direct sum is indexed by the set of nonzero prime ideals of R .

Proof. Let F be a set of r distinct primes p_1, \dots, p_r (i.e., the prime ideals $(p_1), \dots, (p_r)$ are pairwise distinct sets). Let $(m_1, \dots, m_r) \in \prod_{(p) \in F} M_{(p)}$. Then as per the notation and observations section above, T is defined such that

$$T(m_1, \dots, m_r) = m_1 + \dots + m_r$$

We first prove that T is injective. Let $(m_1, \dots, m_r) \in \ker(T)$ be arbitrary. Then $T(m_1, \dots, m_r) = m_1 + \dots + m_r = 0$. By hypothesis, there exist k_1, \dots, k_r such that $p_i^{k_i} m_i = 0$ ($i = 1, \dots, r$). Define $a = p_2^{k_2} \dots p_r^{k_r}$. It follows that $am_2 = \dots = am_r = 0$. Thus,

$$\begin{aligned} a(0) &= 0 \\ a(m_1 + \dots + m_r) &= 0 \\ am_1 + \dots + am_r &= 0 \\ am_1 &= -(am_2 + \dots + am_r) \\ &= -(0 + \dots + 0) \\ &= 0 \end{aligned}$$

Additionally, $\gcd(a, p_1^{k_1}) = 1$ by definition, so $1 \in (a, p_1^{k_1})$. It follows that there exist $b, c \in R$ such that $ba + cp_1^{k_1} = 1$. This combined with the facts that $am_1 = 0$ and $p_1^{k_1} m_1 = 0$ implies that

$$m_1 = 1 \cdot m_1 = (ba + cp_1^{k_1})m_1 = b(am_1) + c(p_1^{k_1} m_1) = b(0) + c(0) = 0$$

A symmetric argument shows that all $m_i = 0$, i.e., $(m_1, \dots, m_r) = (0, \dots, 0)$. Therefore, $\ker(T) = 0$, as desired.

We now prove that T is surjective. Let $m \in \text{Tor}(M)$ be arbitrary. Consider the submodule $N = Am \subset M$. To prove that m is the sum of elements, each from a p -primary component of M , it will suffice to prove that stronger condition that every element in N is the sum of elements, each from a p -primary component of M . Equivalently, it will suffice to show that N is isomorphic to the sum of its p -primary components, since the p -primary components of N are contained in those of M . Define $I = \{a \in R : am = 0\}$. Notice that $I = \ker(l_a)$, where $l_a : R \rightarrow N$ is the left multiplication homomorphism. It follows by the FIT that there exists an isomorphism $\bar{l}_a : R/I \rightarrow N$. Thus, we need only show that R/I is isomorphic to the direct sum of its p -primary components. But the Chinese Remainder Theorem takes care of this for us since I is a nonzero ideal. \square

- In view of the last proposition, our final task will be to classify finitely generated p -primary modules.
- We begin with some definitions.
- **p -primary** (module): An R -module M such that $M = M_{(p)}$ for some prime $p \in R$.
- **Annihilator** (of a module): The set of all $a \in R$ such that $am = 0$ for all $m \in M$. Denoted by $\text{Ann}(M)$. Given by

$$\text{Ann}(M) = \{a \in R : am = 0 \ \forall m \in M\}$$

- **Annihilator** (of an element): The set of all $a \in R$ such that $am = 0$ pertaining to a specific $m \in M$. Denoted by $\text{Ann}(m)$. Given by

$$\text{Ann}(m) = \{a \in R : am = 0\}$$

- Consider $l_m : R \rightarrow M$ defined by $l_m(a) = am$.
 - By the FIT, there exists a module isomorphism $\bar{l}_m : R/\text{Ann}(m) \rightarrow Rm$.

- $\ker(l_m) = \text{Ann}(m)$.
- **Cyclic (module):** An R -module M for which there exists $m \in M$ such that $M = Rm$.
 - Cyclic modules are isomorphic to $R/\text{Ann}(m)$ for a similar reason to the above ($Rm = M$ here).
- With these definitions out of the way, we seek to show that every finitely generated R -module is the direct sum of cyclic modules.
- To prove this result, we will need the following lemma.
- Lemma: Let $M' = Re$ be a cyclic submodule of M . We assume that...
 - (i) $\text{Ann}(e) = (p^n)$;
 - (ii) $p^n M = 0$.

Then every $v \in M/M'$ has a lift $w \in M$ such that $\text{Ann}(w) = \text{Ann}(v)$.

Proof. Let $v \in M/M'$ be arbitrary. Since $p^n M = 0$, $p^n(M/M') = 0$ and hence $\text{Ann}(v) = (p^k)$ for some $k \leq n$. Now let $w \in M$ be an arbitrary lift of v . We will prove that this w satisfies all necessary constraints.

To prove that $\text{Ann}(w) \subset \text{Ann}(v)$, let $a \in \text{Ann}(w)$ be arbitrary. Then $aw = 0$. It follows that $0 = \pi(aw) = a\pi(w) = av$. Therefore, $a \in \text{Ann}(v)$ as well.

To prove that $\text{Ann}(v) \subset \text{Ann}(w)$ □

- Proposition: For every finitely generated p -primary module M , there exist e_1, \dots, e_s such that M is the direct sum of the cyclic submodules Re_i .

Proof. Since M is finitely generated, we know that $M = Rv_1 + \dots + Rv_r$. We induct on r .

For the base case $r = 1$, M is cyclic by definition.

Now suppose that we have proven the claim for some lower cases. Again with the (p^n) issue. □

7.5 Rational Canonical Form and Proofs of Earlier Lemmas

- 2/17:
- Theorem: Every finitely generated R -module M (where R is a PID) is isomorphic to $\text{Tor}(M) \oplus R^h$ for some $h \in \mathbb{Z}_{\geq 0}$, where $h = \text{rank}(M)$.
 - Recall the following theorem.
 - Theorem: Let R be a PID. Then
 - (1) Every finitely generated p -primary R -module is a finite direct sum of cyclic modules (which are isomorphic to $R/p^h R$ for some $h \in \mathbb{N}$).
 - (2) Every torsion module M is the direct sum of its p -primary components.
 - Corollary: Every finitely generated torsion R -module is isomorphic to the finite direct sum of cyclic p -primary modules where p is an element of a finite set of primes. *picture*
 - M finitely generated implies that $M_{(p)}$ is finitely generated.
 - Said aloud that only finite primes p satisfy $M_{(p)} \neq 0$.
 - Theorem (Rational canonical form): Let R be a PID. Then every finitely generated R -torsion module is isomorphic to

$$R/(a_1) \oplus \dots \oplus R/(a_\ell)$$

where $a_2 \mid a_1, a_3 \mid a_2, \dots, a_\ell \mid a_{\ell-1}$.

- Observe: The principal ideal (a_1) is exactly the annihilator of M , i.e.,

$$(a_1) = \{\alpha \in R : \alpha m = 0 \ \forall m \in M\}$$

- Later, (a_1) will play the role of a minimal polynomial, and the product will play the role of the characteristic polynomial.

Proof of theorem. Let p_1, \dots, p_ℓ be ?? the set of primes for which $M_{(p)} \neq 0$. Let

$$M_{(p_i)} \cong R/(p_i^{m_{i,1}}) \times R/(p_i^{m_{i,2}}) \times \dots$$

where $m_{i,1} \geq m_{i,2} \geq \dots$ are such that there exists N for which $m_{i,N} = 0$. Then

$$M/(p_j) \cong R/(p_j^{m_{j,1}})^\times \times R/(p_j^{m_{j,2}})^\times$$

Then we apply the Chinese Remainder Theorem. Define

$$a_r = \prod_{i=1}^{\ell} p_i^{m_{i,r}}$$

where $a_{r+1} \mid a_r$ because $m_{i,j}$ is ?? in j . Use the CRT to imply that

$$\prod_{i=1}^{\ell} R/(p_i^{m_{i,r}}) \cong R/(a_r)$$

□

- That concludes torsion modules over PIDs; we now do torsion modules over fields, which should be easier.
- **R -linearly independent** (elements of M): A set of elements $u_1, \dots, u_\ell \in M$ such that the constraints

$$(a_1, \dots, a_\ell) \in R^\ell \quad \sum_{i=1}^{\ell} a_i u_i = 0$$

imply that $(a_1, \dots, a_\ell) = 0$. Equivalently, $H : R^\ell \rightarrow M$ defined by

$$H(a_1, \dots, a_\ell) = \sum_{i=1}^{\ell} a_i u_i$$

is 1-1, i.e., $R^\ell \cong H(M)$.

- Lemma: Let R be an integral domain, and let M be a finitely generated R -module. Then there exists a submodule $M' \subset M$ such that...

- (i) $M' \cong R^h$ for some $h \geq 0$;

Proof. Let $S \subset M$ be a finite generating set. Select $T \subset S$ such that (i) T is linearly independent and (ii) $T \subsetneq W \subset S$ implies that W is *not* linearly independent. In other words, we are picking T to be a maximal linear independence set. Now suppose $|T| = h$ so that $T = \{u_1, \dots, u_h\}$. Then by definition,

$$M' = \sum_{i=1}^h R u_i \cong R^h$$

where the latter isomorphism follows from Proposition 10.5. □

- (ii) There exists a nonzero $a \in R$ such that $aM \subset M'$ (equivalently, $a(M/M') = 0$).

Proof. Pick $w \in S$ such that $w \notin T$. Then since we picked T to be a *maximal* linear independence set, $T \cup \{w\}$ is linearly *dependent*. It follows that there exists a nonzero $(a_1, \dots, a_{h+1}) \in R^{h+1}$ such that

$$a_1 u_1 + \dots + a_h u_h + a_{h+1} w = 0$$

If $a_{h+1} = 0$, then $(a_1, \dots, a_h) \neq 0$ makes $a_1 u_1 + \dots + a_h u_h = 0$, contradicting the assumed linear independence of T . Thus, $a_{h+1} \neq 0$. It follows that

$$a_{h+1} w = - \sum_{i=1}^h a_i u_i \in M'$$

We may repeat this process for any $w \in S - T$ to obtain a nonzero a_w such that $a_w w \in M'$. Additionally, if $w \in T$, take $a_w = 1$. Now define

$$a = \prod_{w \in S} a_w$$

Since R is an integral domain by hypothesis and each a_w in the above product is nonzero, a is nonzero. Moreover, by its construction, $aw \in M'$ for all $w \in S$. Therefore,

$$aM = a \left(\sum_{s \in S} As \right) \subset M'$$

as desired. □

- Note that you can make stronger statements than the above; you'll just have to use Zorn's lemma to do so.
- We now return to PID-land.
- Corollary: Every finitely generated torsion-free module M over a PID R is isomorphic to R^h for some $h \in \mathbb{Z}_{\geq 0}$.

Proof. Apply the lemma to obtain a submodule M' of M such that $M' \cong R^h$ and a nonzero $a \in R$ such that $aM \subset M'$. Consider $H : M \rightarrow M'$ defined by $H(m) = am$. Since H is just left-multiplication, H is an R -module homomorphism. Additionally, since M is torsion free, $am = 0$ iff $m = 0$ so we have $\ker H = 0$. Thus, since H is injective, $M \cong H(M) \subset M' \cong R^h$. Furthermore, since R is a PID, the submodule $H(M)$ of R^h must be isomorphic to R^n for some $0 \leq n \leq h$ by the Theorem from Week 6. It follows by transitivity that $M \cong H(M) \cong R^n$, as desired. □

- Takeaway: The torsion-free part is far easier to handle than the torsion part.
- Theorem: Let M be a finitely generated R -module, where R is a PID. Then...

- (i) $\text{Tor}(M) \oplus R^h \cong M$ for some $h \geq 0$;

Proof. To prove that $\text{Tor}(M) \oplus R^h \cong M$, the second theorem from Lecture 6.3 tells us that it will suffice to show that $M/\text{Tor}(M) \cong R^h$ for some $h \geq 0$. By part (ii) of the lemma from last time (Lecture 7.2), we have that $M/\text{Tor}(M)$ is torsion-free. This combined with the fact that $M/\text{Tor}(M)$ is a finitely generated (since M is finitely generated) module over a PID allows us to invoke the above corollary, yielding the desired result.

Note that the isomorphism $T : \text{Tor}(M) \oplus R^h \rightarrow M$ is given by

$$T(m, (a_1, \dots, a_h)) = m + \sum a_i e_i$$

where e_1, \dots, e_h generate R^h . □

- (ii) $\text{Tor}(M)$ is finitely generated.

Proof. Since M is finitely generated, part (i) implies that $\text{Tor}(M) \oplus R^h$ is finitely generated. Now consider the projection $\pi : \text{Tor}(M) \oplus R^h \rightarrow \text{Tor}(M)$. Since it is a surjection, the (finite number of) images of the generators of $\text{Tor}(M) \oplus R^h$ generate $\text{Tor}(M)$. \square

- Nori reproves the claim that $M/\text{Tor}(M)$ is torsion-free (see the first lemma from last lecture).
- If $\pi : M \rightarrow M/M'$ and $S : M/M' \rightarrow R^h$ is an isomorphism, then there exists $\varphi : R^h \rightarrow M$ such that the diagram commutes, i.e., $S\pi\varphi = \text{id}_{R^h}$.
- Next week is going to be straight linear algebra.
- Nori would try to do tensors in one week (the last week), but it'd be ridiculous to do something on Friday and put it on a test on Tuesday.
- Imaginary quadratic fields, curves, Dedekind domains, etc.
- Content from this week in the book.
 - Section 12.1.
 - The material before Theorem ?? is OMITTED from the course.
 - Theorem ?? is also OMITTED from the course.
 - The rest of this section will be covered.
 - The main theorems are: The existence theorem (Theorem ??) and the uniqueness theorem (Theorem ??)
 - Section 12.2 deals with the PID $F[X]$ and its applications to linear algebra; this will be covered on Monday next week.

7.6 Office Hours (Callum)

- Problem 6.5?
 - Go with the explicit route, not the universal property of the ring of fractions route.
 - Explicit: Define

$$F(v) = \frac{1}{a}f(av)$$
 - We need to prove that $1/af(av) = 1/bf(bv)$ for valid a, b . Multiply both sides by ab and use commutativity. Thus, $F(v)$ is well defined.
- Problem 6.8?
 - The hardest one. Doesn't really use any of the previous parts.
 - Define $\phi : A \oplus M \rightarrow A^2$ to be the isomorphism. Consider $(1, 0) \in A \oplus M$. In particular, let $\phi(1, 0) = (a, b)$. We know that it will generate a copy of A in A^2 . Essentially, $A(a, b) = A^2$. We know that $\phi^{-1} : A^2 \rightarrow A \oplus M$ and $P : A \oplus M \rightarrow A$. Suppose $P \circ \phi^{-1} : (1, 0) \mapsto c$ and $(0, 1) \mapsto d$.
 - Consider

$$A \hookrightarrow A \oplus M \xrightarrow{\phi} A^2 \xrightarrow{\phi^{-1}} A \oplus M \xrightarrow{P} A$$

which is the identity on A . Then

$$1 \mapsto (1, 0) \mapsto (a, b) = a(1, 0) + b(0, 1) \mapsto ac + bd$$

so $ac + bd = 1$.

- Consider the matrix

$$\begin{pmatrix} a & d \\ b & c \end{pmatrix}$$

- Determinant??

- $(-d, c)$

- So thus, $M = A(-d, c)$??

- $(-d, c) \in A^2$ defines a map from $A^2 \rightarrow M$ with kernel A . $(-d, c) \in \ker(P \circ \phi^{-1})$. Thus, $\phi^{-1}(-d, c) \in \{0\} \oplus M \cong M$.
- Thus, at this point, we may define a map

$$A \hookrightarrow A^2 \xrightarrow{\phi^{-1}} A \oplus M \xrightarrow{P} M$$

by

$$1 \mapsto (-d, c)$$

and this should be an isomorphism.

- $(-d, c)$ generates a submodule of A^2 that is isomorphic to M .
- Injectivity follows from that of all of the components.
- Surjectivity: Pull m back to $(0, m)$ and then $\phi(0, m) \in A^2$. The subset of A^2 equal to all $\phi(0, m)$ is equal to

$$\{(u, v) \in A^2 : \phi^{-1}(u, v) \in 0 \oplus M\} = \{(u, v) \in A^2 : uc + vd = 0\}$$

- We want to find $k \in A$ such that $(u, v) = k(-d, c)$. In other words, we want $u = -kd$ and $v = kc$. $ua = -kda = k(1 - bc) = k - kbc = k - bv$. Thus, $k = ua + bv$. Now we have to substitute that back in and show that it works.
- Thus, we have that

$$kc = ua + bvc = uac + b(1 - ad) = v + uac - vad = v + a(bc - ad)$$

- Saying $A \cong M$ is kind of like saying that there's a change of basis. That's why matrices keep coming up.
- Summary of what we did.

1. We have

$$A \hookrightarrow A \oplus M \xrightarrow{\phi} A^2 \xrightarrow{\phi^{-1}} A \oplus M \xrightarrow{P} A$$

and this is the identity.

2. We define $(1, 0) \mapsto (a, b)$, which will generate a copy of A in A^2 .
3. We now need to find a basis vector corresponding to M (which we hope is A).
4. $\{(1, 0), (0, 1)\}$ is the standard basis for A^2 .
5. We need to solve for x, y such that

$$\begin{pmatrix} a & x \\ b & y \end{pmatrix}$$

is invertible.

6. $\{\phi^{-1}(1, 0), \phi^{-1}(0, 1)\}$ is another basis of A^2 .
7. We want $ac + bd = 1$.

7.7 Chapter 11: Vector Spaces

From Dummit and Foote (2004).

Section 11.1: Definitions and Basic Theory

- 2/20:
- Reviewing Labalme (2021) is probably a good idea.
 - Many of Dummit and Foote (2004)'s proofs more elegant, though.
 - Goal of this chapter:
 - Brief overview of results that will be used later on; more in-depth (even introductory level) linear algebra topics, such as Gauss-Jordan elimination, row echelon forms, etc., will not be covered.
 - Only finite-dimensional vector spaces are discussed in the text; some stuff on infinite dimensional vector spaces is included in the exercises
 - Characteristic polynomials and eigenvalues: Next chapter.
 - Module terminology vs. vector space terminology.

Terminology for R any Ring	Terminology for R a Field
M is an R -module	M is a vector space over R
m is an element of M	m is a vector in M
α is a ring element	α is a scalar
N is a submodule of M	N is a subspace of M
M/N is a quotient module	M/N is a quotient space
M is a free module of rank n	M is a vector space of dimension n
M is a finitely generated module	M is a finite dimensional vector space
M is a nonzero cyclic module	M is a 1-dimensional vector space
$\varphi : M \rightarrow N$ is an R -module homomorphism	$\varphi : M \rightarrow N$ is a linear transformation
M and N are isomorphic as R -modules	M and N are isomorphic vector spaces
The subset A of M generates M	The subset A of M spans M
$M = RA$	Each element of M is a linear combination of elements of A , i.e., $M = \text{Span}(A)$

Table 7.1: Module vs. vector space terminology.

- In this chapter, F denotes a field and V denotes a vector space over F .
- **Linearly independent** (subset $S \subset V$): A subset S of V for which the equation $\alpha_1 v_1 + \cdots + \alpha_n v_n = 0$ with $\alpha_1, \dots, \alpha_n \in F$ and $v_1, \dots, v_n \in S$ implies $\alpha_1 = \cdots = \alpha_n = 0$.
- **Basis**: An ordered set of linearly independent vectors which span V . *Also known as ordered basis.*
 - In particular, two bases will be considered different even if one is simply a rearrangement of the other.
- Examples.
 1. $V = F[X]$.
 - Basis: $1, X, X^2, \dots$ is linearly independent by definition since a polynomial is zero iff all of its coefficients are 0.
 2. The collection of solutions of a linear, homogeneous, constant coefficient differential equation over \mathbb{C} .
 - A vector space since differentiation is a linear operator.
 - Elements are linearly independent if they are linearly independent as functions.
 - Example: e^t, e^{2t} are easily seen to be solutions of the equation $y'' - 3y' + 2y = 0$.

- They are linearly independent since $ae^t + be^{2t} = 0$ implies $a + b = 0$ ($t = 0$) and $ae + be^2 = 0$ ($t = 1$), and the only solution to this system of two equations is $a = b = 0$.
- It is a theorem of differential equations that these elements span the set of solutions of this equation.

- Vector spaces are free modules.

Proposition 11.1. Assume the set $\mathcal{A} = \{v_1, \dots, v_n\}$ spans the vector space V but no proper subset of \mathcal{A} spans V . Then \mathcal{A} is a basis of V . In particular, any finitely generated (i.e., finitely spanned) vector space over F is a free F -module.

Proof. Given. □

- Example.

1. Consider $F[X]/(f)$, where $f = X^n + a_{n-1}X^{n-1} + \dots + a_0$.
 - (f) is a subspace of $F[X]$.
 - Euclidean Algorithm: Every $a \in F[X]$ can be written uniquely in the form $qf + r$ where $0 \leq \deg(r) \leq n - 1$. Thus, every element of the quotient is represented by a polynomial r of degree $\leq n - 1$.
 - It follows that $\overline{1}, \overline{X}, \overline{X^2}, \dots, \overline{X^{n-1}}$ spans $F[X]/(f)$.

- Spanning sets contain bases.

Corollary 11.2. Assume the finite set \mathcal{A} spans the vector space V . Then \mathcal{A} contains a basis of V .

Proof. Given. □

- A new property of bases.

Theorem 11.3 (Replacement Theorem). Assume $\mathcal{A} = \{a_1, \dots, a_n\}$ is a basis for V containing n elements and $\{b_1, \dots, b_m\}$ is a set of linearly independent vectors in V . Then there is an ordering a_1, \dots, a_n such that for each $k \in \{1, \dots, m\}$, the set

$$\{b_1, \dots, b_k, a_{k+1}, \dots, a_n\}$$

is a basis of V . In other words, the elements b_1, \dots, b_m can be used to successively replace the elements of the basis \mathcal{A} , still retaining a basis. In particular, $n \geq m$.

Proof. Given. □

- Linear independence, span, and cardinality.

Corollary 11.4.

1. Suppose V has a finite basis with n elements. Any set of linearly independent vectors has $\leq n$ elements. Any spanning set has $\geq n$ elements.
2. If V has some finite basis, then any two bases of V have the same cardinality.

Proof. Given. □

- **Dimension:** The cardinality of any basis of V . Denoted by $\dim_F V$, $\dim V$.
- **Finite dimensional** (vector space): A vector space V that is finitely generated.
- **Infinite dimensional** (vector space): A vector space V that is not finitely generated.

– We write $\dim V = \infty$ for these.

- Examples.

1. The dimension of the solution space to $y'' - 3y' + 2y = 0$ is 2.

– Recall from above that a basis is e^t, e^{2t} .

– In general, it is a theorem in differential equations that the space of solutions of an n^{th} order linear, homogeneous, constant coefficient differential equation of degree n over \mathbb{C} is a vector space over \mathbb{C} of dimension n .

2. The dimension of $F[X]/(f)$ is $\deg(f)$.

– $F[X]$ and (f) are infinite dimensional vector spaces.

- Linearly independent lists and bases.

Corollary 11.5 (Building-Up Lemma). If A is a set of linearly independent vectors in the finite dimensional space V , then there exists a basis of V containing A .

Proof. Given. □

- Characterizing finite dimensional vector spaces.

Theorem 11.6. If V is an n -dimensional vector space over F , then $V \cong F^n$. In particular, any two finite dimensional vector spaces over F of the same dimension are isomorphic.

Proof. Given. □

- Examples.

1. Bases of \mathbb{F}_q^k .

– Dummit and Foote (2004) justifies that the number of distinct bases of \mathbb{F}_q^k is

$$(q^k - 1)(q^k - q)(q^k - q^2) \cdots (q^k - q^{k-1})$$

– For every vector $v \in \mathbb{F}_q^k$, there are $q - 1$ other linearly dependent vectors (corresponding to the q \mathbb{F} -multiples of it).

2. Subspaces of \mathbb{F}_q^n .

– Dummit and Foote (2004) justifies that the number of distinct k -dimensional subspaces of \mathbb{F}_q^n is

$$\frac{(q^n - 1)(q^n - q) \cdots (q^n - q^{k-1})}{(q^k - 1)(q^k - q) \cdots (q^k - q^{k-1})}$$

- Dimension of the quotient space.

Theorem 11.7. Let V be a vector space over F , and let W be a subspace of V . Then V/W is a vector space with $\dim V = \dim W + \dim V/W$ (where if one side is infinite, then both are).

Proof. Given. □

- Dimension of the kernel and image of a linear transformation.

Corollary 11.8. Let $\varphi : V \rightarrow U$ be a linear transformation of vector spaces over F . Then $\ker \varphi$ is a subspace of V , $\varphi(V)$ is a subspace of U , and $\dim V = \dim \ker \varphi + \dim \varphi(V)$.

Proof. Given. □

- Classifying isomorphic operator.

Corollary 11.9. Let $\varphi : V \rightarrow W$ be a linear transformation of vector spaces of the same finite dimension. Then the following are equivalent.

1. φ is an isomorphism.
2. φ is injective, i.e., $\ker \varphi = 0$.
3. φ is surjective, i.e., $\varphi(V) = W$.
4. φ sends a basis of V to a basis of W .

Proof. Given. □

- **Null space** (of a linear transformation): The kernel of the linear transformation.
- **Nullity** (of a linear transformation): The dimension of the kernel of the linear transformation.
- **Rank** (of a linear transformation): The dimension of the image of the linear transformation.
- **Nonsingular** (linear transformation): A linear transformation φ for which $\ker \varphi = 0$.
- **General linear group**: The group of all nonsingular linear transformations from $V \rightarrow V$ under the group operation of composition. Denoted by $GL(V)$.

– Dummit and Foote (2004) justifies that if $V = \mathbb{F}_q^n$, then

$$|GL(V)| = (q^n - 1)(q^n - q) \cdots (q^n - q^{n-1})$$

Exercises

4. Prove that the space of real-valued functions on the closed interval $[a, b]$ is an infinite dimensional vector space over \mathbb{R} , where $a < b$.
5. Prove that the space of continuous real-valued functions on the closed interval $[a, b]$ is an infinite dimensional vector space over \mathbb{R} , where $a < b$.
10. Prove that any vector space V has a basis (by convention, the null set is the basis for the zero space). *Hint:* Let \mathcal{S} be the set of subsets of V consisting of linearly independent vectors, partially ordered under inclusion; apply Zorn's Lemma to \mathcal{S} and show that a maximal element of \mathcal{S} is a basis.
11. Refine your argument in the preceding exercise to prove that any set of linearly independent vectors of V is contained in a basis of V .
12. If F is a field with a finite or countable number of elements and V is an infinite dimensional vector space over F with basis \mathcal{B} , prove that the cardinality of V equals the cardinality of \mathcal{B} . Deduce in this case that any two bases of V have the same cardinality.
13. Prove that as vector spaces over \mathbb{Q} , $\mathbb{R}^n \cong \mathbb{R}$ for all $n \in \mathbb{Z}^+$. Note that, in particular, this means that \mathbb{R}^n and \mathbb{R} are isomorphic as additive abelian groups.
14. Let \mathcal{A} be a basis for the infinite dimensional vector space V . Prove that V is isomorphic to the direct sum of copies of the field F indexed by the set \mathcal{A} . Prove that the direct product of copies of F indexed by \mathcal{A} is a vector space over F and it has strictly larger dimension than the dimension of V (see the exercises in Section 10.3 for the definitions of direct sum and direct product over infinitely many modules).

Section 11.2: The Matrix of a Linear Transformation

- Assumptions for this section.
 - V, W are vector spaces over the field F .
 - $\mathcal{B} = \{v_1, \dots, v_n\}$ is an (ordered) basis of V , and $\mathcal{E} = \{w_1, \dots, w_m\}$ is an (ordered) basis of W .
 - $\varphi \in \text{Hom}(V, W)$.

- Matrix** (of φ with respect to the bases \mathcal{B}, \mathcal{E}): The $m \times n$ matrix whose i, j entry is α_{ij} , where

$$\varphi(v_j) = \sum_{i=1}^m \alpha_{ij} w_i$$

Denoted by $M_{\mathcal{B}}^{\mathcal{E}}(\varphi)$.

- Dummit and Foote (2004) reviews how to recover φ from $M_{\mathcal{B}}^{\mathcal{E}}(\varphi)$.
 - The equivalence of matrix multiplying and linear transforming is sometimes denoted

$$[\varphi(v)]_{\mathcal{E}} = M_{\mathcal{B}}^{\mathcal{E}}(\varphi)[v]_{\mathcal{B}}$$

- Representation** (of φ with respect to the bases \mathcal{B}, \mathcal{E}): The matrix $A = (a_{ij})$ associated with φ .
- Examples.

- Computing a matrix with respect to the standard bases of $\mathbb{R}^3, \mathbb{R}^2$.
- The matrix of the differentiation operator $\varphi : V \rightarrow V$ on the 2-dimensional space of solutions V to $y'' - 3y' + 2y = 0$.
 - Since

$$\varphi(v_1) = \frac{d}{dt}(e^t) = e^t = v_1 \qquad \varphi(v_2) = \frac{d}{dt}(e^{2t}) = 2e^{2t} = 2v_2$$

the representation of φ is

$$\begin{pmatrix} 1 & 0 \\ 0 & 2 \end{pmatrix}$$

- Computing a matrix with respect to the standard bases of $\mathbb{Q}^3, \mathbb{Q}^3$.
- Isomorphism between the space of linear transformations and the space of matrices.

Theorem 11.10. Let V be a vector space over F of dimension n and let W be a vector space over F of dimension m , with respective bases \mathcal{B}, \mathcal{E} . Then the map $\text{Hom}_F(V, W) \rightarrow M_{m \times n}(F)$ from the space of linear transformations from V to W to the space of $m \times n$ matrices with coefficients in F defined by $\varphi \mapsto M_{\mathcal{B}}^{\mathcal{E}}(\varphi)$ is a vector space isomorphism. In particular, there is a bijective correspondence between linear transformations and their associated matrices with respect to a fixed choice of bases.

Proof. Given. □

- There is no *natural* isomorphism between $\text{Hom}_F(V, W)$ and $M_{m \times n}(F)$.
 - This is because the choices of bases are arbitrary (there is no natural choice of them).
- Dimension of the space of linear transformations.

Corollary 11.11. The dimension of $\text{Hom}_F(V, W)$ is $(\dim V)(\dim W)$.

Proof. Given. □

- **Nonsingular** (matrix): An $m \times n$ matrix A such that $Ax = 0$ with $x \in F^n$ implies that $x = 0$. Also known as **invertible**.
- Nonsingular linear transformations vs. nonsingular matrices.
 - Independent of the choice of bases, a matrix is nonsingular iff the corresponding linear transformation is nonsingular.
- Dummit and Foote (2004) uses the definition of the matrix to deduce the formula for matrix multiplication.
- Relating matrix multiplication to linear transformation composition.

Theorem 11.12. Let U, V, W be finite dimensional vector spaces over F with ordered bases $\mathcal{D}, \mathcal{B}, \mathcal{E}$, and assume $\psi : U \rightarrow V$ and $\varphi : V \rightarrow W$ are linear transformations. Then

$$M_{\mathcal{D}}^{\mathcal{E}}(\varphi \circ \psi) = M_{\mathcal{B}}^{\mathcal{E}}(\varphi) M_{\mathcal{D}}^{\mathcal{B}}(\psi)$$

In words, the product of the matrices representing the linear transformations φ, ψ is the matrix representing the composite linear transformation $\varphi \circ \psi$.

- Properties of matrix multiplication.

Corollary 11.13. Matrix multiplication is associative and distributive (whenever the dimensions are such as to make products defined). An $n \times m$ matrix A is nonsingular if and only if it is invertible.

Proof. Given. □

- Ring-like properties of $M_n(F)$, as induced by those of $\text{Hom}_F(V, V)$.

Corollary 11.14.

1. If \mathcal{B} is a basis of the n -dimensional space V , the map $\varphi \mapsto M_{\mathcal{B}}^{\mathcal{B}}(\varphi)$ is a ring and a vector space isomorphism of $\text{Hom}_F(V, V)$ onto the space $M_n(F)$ of $n \times n$ matrices with coefficients in F .
2. $GL(V) \cong GL_n(F)$, where $\dim V = n$. In particular, if F is a finite field, the order of the finite group $GL_n(F)$ (which equals $|GL(V)|$) is given by the formula at the end of Section 11.1.

Proof. Given. □

- **Row rank** (of a matrix): The maximal number of linearly independent rows of the matrix, where the rows are considered as vectors in affine m -space.
- **Column rank** (of a matrix): The maximal number of linearly independent columns of the matrix, where the columns are considered as vectors in affine n -space.
- Relating ranks.
 - The rank of ψ equals the column rank of $M_{\mathcal{B}}^{\mathcal{E}}(\psi)$.
- **Similar** (matrices): Two $n \times n$ matrices A, B for which there exists an invertible $n \times n$ matrix P such that $P^{-1}AP = B$.
- **Similar** (linear transformations): Two linear transformations $\varphi, \psi : V \rightarrow V$ for which there exists a nonsingular linear transformation ξ such that $\xi^{-1}\varphi\xi = \psi$.
 - This is an equivalence relation whose equivalence classes are the orbits of $GL(V)$ acting by conjugation on $\text{Hom}_F(V, V)$.

- **Transition** (matrix from \mathcal{B} to \mathcal{E}): The matrix defined as follows, where I is the identity transformation. Also known as **change of basis** (matrix). Denoted by P . Given by

$$P = M_{\mathcal{B}}^{\mathcal{E}}(I)$$

- $P = M_{\mathcal{B}}^{\mathcal{E}}(I)$ satisfies $P^{-1}M_{\mathcal{B}}^{\mathcal{B}}(I)P = M_{\mathcal{E}}^{\mathcal{E}}(\varphi)$.
 - If $\mathcal{B} \neq \mathcal{E}$, then P is not the identity matrix.
- Note that we need *ordered* bases to have a unique $P = M_{\mathcal{B}}^{\mathcal{E}}(I)$!
- **Change of basis**: The similarity action of $M_{\mathcal{B}}^{\mathcal{E}}(I)$ on $M_{\mathcal{B}}^{\mathcal{B}}(\varphi)$.
- Dummit and Foote (2004) proves that any two similar matrices represent the same linear transformation with respect to two different choices of bases.
- Example of similarity given.
- **Canonical forms**: The study of the simplest possible matrix representing a given linear transformation (and which basis to choose to realize it).
- We now move on to linear transformations on tensor products of vector spaces.
- Return to later.
- **Idempotent** (linear transformation): A linear transformation ψ satisfying $\psi^2 = \psi$.
 - Characterized in Exercise 11.2.11.

Section 11.3: Dual Vector Spaces

- **Dual space** (of a vector space): The space of linear transformations from V to F . Denoted by V^* .
- **Linear functional**: An element of V^* .
- **Dual basis** (to a basis of V): The basis related to a basis $\{v_1, \dots, v_n\}$ of V by

$$v_i^*(v_j) = \begin{cases} 1 & i = j \\ 0 & i \neq j \end{cases}$$

for $1 \leq j \leq n$. Denoted by $\{v_1^*, \dots, v_n^*\}$.

- The dual basis to a basis of V is a basis of V^* .

Proposition 11.18. With notations as above, $\{v_1^*, \dots, v_n^*\}$ is a basis of V^* . In particular, if V is finite dimensional, then V^* has the same dimension as V .

Proof. Given. □

- If V is infinite dimensional, then $\dim V < \dim V^*$.
- **Algebraic** (dual space to V): The dual space V^* taken for V of arbitrary dimension.
- If V has additional structure (e.g., a topology), we can get other types of dual spaces, such as the following.
- **Continuous** (dual of V): A dual of V in which the linear functionals must be continuous.
- Example.
 1. Let $V = C([a, b], \mathbb{R})$.
 - If $a < b$, then V is infinite dimensional.

- For each $g \in V$, the function $\varphi_g : V \rightarrow \mathbb{R}$ defined by

$$\varphi_g(f) = \int_a^b f(t)g(t) dt$$

is a linear functional on V .

- **Double dual** (of V): The dual of V^* . Also known as **second dual**. Denoted by V^{**} .
- For finite dimensional V , $\dim V = \dim V^{**}$ and hence $V \cong V^{**}$.
 - There is a **natural** (i.e., basis independent/coordinate free) isomorphism.
 - More detail on this is given.
 - This is different for infinite dimensional V , as per the above.

- Existence of a natural map $V \rightarrow V^{**}$.

Theorem 11.19. There is a natural injective linear transformation from V to V^{**} . If V is finite dimensional, then this linear transformation is an isomorphism.

Proof. Given. □

- φ^* : The induced function from $W^* \rightarrow V^*$ defined by

$$f \mapsto f \circ \varphi$$

- This is just the **pullback** or **dual map**.

- Pullback: Linearity and matrix.

Theorem 11.20. With notations as above, φ^* is a linear transformation from W^* to V^* and $M_{\mathcal{E}^*}^{\mathcal{B}^*}(\varphi^*)$ is the transpose of the matrix $M_{\mathcal{B}}^{\mathcal{E}}(\varphi)$.

Proof. Given. □

- A partial statement of the rank-nullity theorem.

Corollary 11.21. For any matrix A , the row rank of A equals the column rank of A .

Proof. Given. □

- **Annihilator** (of S in V): The set of all $v \in V$ for which $f(v) = 0$ for all $f \in S \subset V^*$. Denoted by $\text{Ann}(S)$. Given by

$$\text{Ann}(S) = \{v \in V : f(v) = 0 \forall f \in S\}$$